


Keras

 Star

52,882

About Keras

Getting started

Developer guides

Keras API reference

Code examples

Computer Vision

Natural Language Processing

Structured Data

Timeseries

Audio Data

Generative Deep Learning

Reinforcement Learning

Graph Data

Quick Keras Recipes

Why choose Keras?

Community & governance

Contributing to Keras

KerasTuner

» [Code examples](#) / [Natural Language Processing](#) / Bidirectional LSTM on IMDB

Bidirectional LSTM on IMDB

Author: [fchollet](#)
Date created: 2020/05/03
Last modified: 2020/05/03
Description: Train a 2-layer bidirectional LSTM on the IMDB movie review sentiment classification dataset.

 [View in Colab](#) •  [GitHub source](#)

Setup

```
import numpy as np
from tensorflow import keras
from tensorflow.keras import layers

max_features = 20000 # Only consider the top 20k words
maxlen = 200 # Only consider the first 200 words of each movie review
```

Build the model

```
# Input for variable-length sequences of integers
inputs = keras.Input(shape=(None,), dtype="int32")
# Embed each integer in a 128-dimensional vector
x = layers.Embedding(max_features, 128)(inputs)
# Add 2 bidirectional LSTMs
x = layers.Bidirectional(layers.LSTM(64, return_sequences=True))(x)
x = layers.Bidirectional(layers.LSTM(64))(x)
# Add a classifier
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.summary()
```

Model: "model"		
Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, None)]	0
=====		
embedding (Embedding)	(None, None, 128)	2560000
=====		
bidirectional (Bidirectional	(None, None, 128)	98816
=====		
bidirectional_1 (Bidirection	(None, 128)	98816
=====		
dense (Dense)	(None, 1)	129
=====		
Total params: 2,757,761		
Trainable params: 2,757,761		
Non-trainable params: 0		

Load the IMDB movie review sentiment data

```
(x_train, y_train), (x_val, y_val) = keras.datasets.imdb.load_data(
    num_words=max_features
)
print(len(x_train), "Training sequences")
print(len(x_val), "Validation sequences")
x_train = keras.preprocessing.sequence.pad_sequences(x_train, maxlen=maxlen)
x_val = keras.preprocessing.sequence.pad_sequences(x_val, maxlen=maxlen)
```

25000 Training sequences
25000 Validation sequences

Train and evaluate the model

```
model.compile("adam", "binary_crossentropy", metrics=["accuracy"])
model.fit(x_train, y_train, batch_size=32, epochs=2, validation_data=(x_val, y_val))
```

Epoch 1/2
782/782 [=====] - 220s 281ms/step - loss: 0.4117 - accuracy: 0.8083
- val_loss: 0.6497 - val_accuracy: 0.6983
Epoch 2/2
726/782 [=====>...] - ETA: 11s - loss: 0.3170 - accuracy: 0.8683